Homework 1

Zach Lees

Table of contents

T : 1 / /1 O:	. 1	1				٠,																										
Appendix															٥																	
Question 3	•	•	•				•	•	•	٠	•	•	•	•		•	•	•	٠			•		•		•	•		•		•	4
Question 2																																٠
Question 1																																2

Link to the Github repository

Due: Sun, Jan 29, 2023 @ 11:59pm

Please read the instructions carefully before submitting your assignment.

- 1. This assignment requires you to:
 - Upload your Quarto markdown files to a git repository
 - Upload a PDF file on Canvas
- 2. Don't collapse any code cells before submitting.
- 3. Remember to make sure all your code output is rendered properly before uploading your submission.

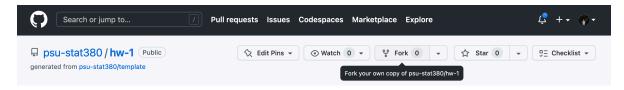
Please add your name to the the author information in the frontmatter before submitting your assignment.

Question 1



In this question, we will walk through the process of *forking* a git repository and submitting a *pull request*.

1. Navigate to the Github repository here and fork it by clicking on the icon in the top right



Provide a sensible name for your forked repository when prompted.

2. Clone your Github repository on your local machine

```
$ git clone <<https://github.com/ZacharyLees/hw-1.git>>
$ cd hw-1
```

Alternatively, you can use Github codespaces to get started from your repository directly.

3. In order to activate the R environment for the homework, make sure you have renv installed beforehand. To activate the renv environment for this assignment, open an instance of the R console from within the directory and type

```
renv::activate()
```

Follow the instrutions in order to make sure that renv is configured correctly.

- 4. Work on the *reminaing part* of this assignment as a .qmd file.
 - Create a PDF and HTML file for your output by modifying the YAML frontmatter for the Quarto .qmd document
- 5. When you're done working on your assignment, push the changes to your github repository.
- 6. Navigate to the original Github repository here and submit a pull request linking to your repository.

Remember to include your name in the pull request information!

If you're stuck at any step along the way, you can refer to the official Github docs here

Question 2



Consider the following vector

```
my_vec <- c(
    "+0.07",
    "-0.07",
    "+0.25",
    "-0.84",
    "+0.32",
    "-0.24",
    "-0.97",
    "-0.36",
    "+1.76",
    "-0.36")
```

For the following questions, provide your answers in a code cell.

1. What data type does the vector contain?

#chr

1. Create two new vectors called my_vec_double and my_vec_int which converts my_vec to Double & Integer types, respectively,

```
#converts the vector to an int
my_vec_int <- as.integer(my_vec)

#converts the vector to a double, and checks is it's classified as a double
my_vec_double <- as.double(my_vec)
is.double(my_vec_double)</pre>
```

[1] TRUE

- 1. Create a new vector my_vec_bool which comprises of:
 - TRUEif an element in my_vec_double is ≤ 0
 - FALSE if an element in my_vec_double is ≥ 0

#returns true if the number is greater than 0, and false if the number is less than 0 my_vec_bool <- $(my_vec_double > 0)$

How many elements of `my_vec_double` are greater than zero?

```
#sums all the numbers greater than 0
sum(my_vec_bool)
```

[1] 4

1. Sort the values of my_vec_double in ascending order.

```
sort(my_vec_double)
```

Question 3



In this question we will get a better understanding of how R handles large data structures in memory.

1. Provide R code to construct the following matrices:

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} \quad \text{and} \quad \begin{bmatrix} 1 & 2 & 3 & 4 & 5 & \dots & 100 \\ 1 & 4 & 9 & 16 & 25 & \dots & 10000 \end{bmatrix}$$

```
A Tip
```

Recall the discussion in class on how R fills in matrices

In the next part, we will discover how knowledge of the way in which a matrix is stored in memory can inform better code choices. To this end, the following function takes an input n and creates an $n \times n$ matrix with random entries.

For example:

```
generate_matrix(4)
```

```
[,1] [,2] [,3] [,4] [1,] [1,] 0.3729221 -0.7468015 -0.5402618 0.17509000 [2,] -1.1947182 -0.2584650 1.4023170 0.92446498 [3,] 0.4202974 -0.6709274 0.7698156 -0.05129285 [4,] -0.1010081 1.1765770 -2.0807964 0.66939293
```

Let M be a fixed 50×50 matrix

```
M <- generate_matrix(50)
mean(M)</pre>
```

[1] 0.02320708

2. Write a function row_wise_scan which scans the entries of M one row after another and outputs the number of elements whose value is ≥ 0 . You can use the following starter code

```
#makes a function that allows us to determine amount of positive in the matrix
row_wise_scan <- function(x){
    n <- nrow(x)</pre>
```

```
m <- ncol(x)

# Insert your code here
count <- 0
for(n in 1:50){
    for(m in n:50){
        if(m > 0){
            count <- count + 1
        }
    }
}

return(count)
}</pre>
```

[1] 1275

```
#returns false for values less than 0, and true for values greater than 0 scan <- (M > 0) #returns the number of true values sum(scan)
```

[1] 1253

3. Similarly, write a function col_wise_scan which does exactly the same thing but scans the entries of M one column after another

```
col_wise_scan <- function(x){
    count <- 0

    # Insert your code here
    count <- 0
    for(m in 1:50){
        for(n in m:50){
            if(m > 0){
                count <- count + 1
            }
        }</pre>
```

```
return(count)
}
col_wise_scan(M)
```

[1] 1275

You can check if your code is doing what it's supposed to using the function here¹

4. Between col_wise_scan and row_wise_scan, which function do you expect to take shorter to run? Why? ::: {.cell}

```
#col_wise_scan, because R by default processes info by column.
```

:::

5. Write a function time_scan which takes in a method f and a matrix M and outputs the amount of time taken to run f(M) ::: {.cell}

```
time_scan <- function(f, M){
   initial_time <- Sys.time() # Write your code here
   f(M)
   final_time <- Sys.time() # Write your code here

   total_time_taken <- final_time - initial_time
   return(total_time_taken)
}</pre>
```

:::

Provide your output to ::: {.cell}

```
sapply(1:100, function(i) {
    x <- generate_matrix(100)
    row_wise_scan(x) == col_wise_scan(x)
}) %>% sum == 100
```

 $^{^1\}mathrm{If}$ your code is right, the following code should evaluate to be \mathtt{TRUE}

```
list(
      row_wise_time = time_scan(row_wise_scan, M),
       col_wise_time = time_scan(row_wise_scan, M)
  )
$row_wise_time
Time difference of 5.698204e-05 secs
$col_wise_time
Time difference of 5.722046e-05 secs
:::
Which took longer to run? row_wise_time took longer to run.
  6. Repeat this experiment now when:
       • M is a 100 \times 100 matrix ::: {.cell}
  #generates a new matrix as 100 x 100
  M <- generate_matrix(100)</pre>
  #gets times for this new matrix
  list(
       row_wise_time = time_scan(row_wise_scan, M),
       col_wise_time = time_scan(row_wise_scan, M)
  )
$row_wise_time
Time difference of 5.412102e-05 secs
$col_wise_time
Time difference of 5.197525e-05 secs
:::
* `M` is a $1000 \times 1000$ matrix
  M <- generate_matrix(1000)</pre>
  list(
```

```
row_wise_time = time_scan(row_wise_scan, M),
      col_wise_time = time_scan(row_wise_scan, M)
  )
$row_wise_time
Time difference of 5.793571e-05 secs
$col_wise_time
Time difference of 5.102158e-05 secs
* `M` is a $5000 \times 5000$ matrix
  M <- generate_matrix(5000)</pre>
  list(
      row_wise_time = time_scan(row_wise_scan, M),
      col_wise_time = time_scan(row_wise_scan, M)
  )
$row_wise_time
Time difference of 6.103516e-05 secs
$col_wise_time
Time difference of 5.197525e-05 secs
```

What can you conclude? The larger the matrix becomes the longer it takes to run the code on average.

Appendix

Print your R session information using the following command

```
sessionInfo()
```

R version 4.2.2 (2022-10-31 ucrt)

Platform: x86_64-w64-mingw32/x64 (64-bit)
Running under: Windows 10 x64 (build 22621)

Matrix products: default

locale:

- [1] LC_COLLATE=English_United States.utf8
- [2] LC_CTYPE=English_United States.utf8
- [3] LC_MONETARY=English_United States.utf8
- [4] LC_NUMERIC=C
- [5] LC_TIME=English_United States.utf8

attached base packages:

[1] stats graphics grDevices datasets utils methods base

loaded via a namespace (and not attached):

[1] compiler_4.2.2 fastmap_1.1.0 cli_3.6.0 htmltools_0.5.4 [5] tools_4.2.2 rstudioapi_0.14 yaml_2.3.7 rmarkdown_2.20

[9] knitr_1.42 xfun_0.36 digest_0.6.31 jsonlite_1.8.4

[13] rlang_1.0.6 renv_0.16.0-53 evaluate_0.20