

Homework 1

Varun Bhatnagar

Table of contents

Question 1	2
Question 2	3
Question 3	4

Appendix	8
-----------------	----------

[Link to the Github repository](#)

! Due: Sun, Jan 29, 2023 @ 11:59pm

Please read the instructions carefully before submitting your assignment.

1. This assignment requires you to:
 - Upload your Quarto markdown files to a `git` repository
 - Upload a PDF file on Canvas
2. Don't collapse any code cells before submitting.
3. Remember to make sure all your code output is rendered properly before uploading your submission.

Please add your name to the the author information in the frontmatter before submitting your assignment.

Question 1

💡 20 points

In this question, we will walk through the process of *forking* a `git` repository and submitting a *pull request*.

1. Navigate to the Github repository [here](#) and fork it by clicking on the icon in the top right



Provide a sensible name for your forked repository when prompted.

2. Clone your Github repository on your local machine

```
$ git clone <<insert your repository url here>>
$ cd hw-1
```

Alternatively, you can use Github codespaces to get started from your repository directly.

3. In order to activate the R environment for the homework, make sure you have `renv` installed beforehand. To activate the `renv` environment for this assignment, open an instance of the R console from within the directory and type

```
renv::activate()
```

Follow the instructions in order to make sure that `renv` is configured correctly.

4. Work on the *remaining part* of this assignment as a `.qmd` file.
 - Create a PDF and HTML file for your output by modifying the YAML frontmatter for the Quarto `.qmd` document
5. When you're done working on your assignment, push the changes to your github repository.
6. Navigate to the original Github repository [here](#) and submit a pull request linking to your repository.

Remember to **include your name** in the pull request information!

If you're stuck at any step along the way, you can refer to the [official Github docs here](#)

Question 2

💡 30 points

Consider the following vector

```
my_vec <- c(
  "+0.07",
  "-0.07",
  "+0.25",
  "-0.84",
  "+0.32",
  "-0.24",
  "-0.97",
  "-0.36",
  "+1.76",
  "-0.36"
)
```

For the following questions, provide your answers in a code cell.

1. What data type does the vector contain?

```
# String
```

1. Create two new vectors called `my_vec_double` and `my_vec_int` which converts `my_vec` to Double & Integer types, respectively,

```
my_vec_double <- as.double(my_vec)
my_vec_int <- as.integer(my_vec)
```

1. Create a new vector `my_vec_bool` which comprises of:
 - TRUE if an element in `my_vec_double` is ≤ 0
 - FALSE if an element in `my_vec_double` is ≥ 0

```
my_vec_bool <- my_vec_double <= 0
```

How many elements of `my_vec_double` are greater than zero?

```
length(my_vec_bool) - sum(my_vec_bool)
```


```
[1] 4
```

1. Sort the values of `my_vec_double` in ascending order.

```
sort(my_vec_double)
```

```
[1] -0.97 -0.84 -0.36 -0.36 -0.24 -0.07  0.07  0.25  0.32  1.76
```

Question 3

 50 points

In this question we will get a better understanding of how R handles large data structures in memory.

1. Provide R code to construct the following matrices:

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} \quad \text{and} \quad \begin{bmatrix} 1 & 2 & 3 & 4 & 5 & \dots & 100 \\ 1 & 4 & 9 & 16 & 25 & \dots & 10000 \end{bmatrix}$$

::: {.cell}

```
matrix(1:9, nrow = 3)
```

```
      [,1] [,2] [,3]
[1,]     1     4     7
[2,]     2     5     8
[3,]     3     6     9
```

```
t(cbind(c(1:100), c(1:100) ^ 2))
```

```
      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13] [,14]
[1,]     1     2     3     4     5     6     7     8     9    10    11    12    13    14
[2,]     1     4     9    16    25    36    49    64    81   100   121   144   169   196
      [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22] [,23] [,24] [,25] [,26]
[1,]    15    16    17    18    19    20    21    22    23    24    25    26
[2,]   225   256   289   324   361   400   441   484   529   576   625   676
      [,27] [,28] [,29] [,30] [,31] [,32] [,33] [,34] [,35] [,36] [,37] [,38]
[1,]    27    28    29    30    31    32    33    34    35    36    37    38
[2,]   729   784   841   900   961  1024  1089  1156  1225  1296  1369  1444
      [,39] [,40] [,41] [,42] [,43] [,44] [,45] [,46] [,47] [,48] [,49] [,50]
[1,]    39    40    41    42    43    44    45    46    47    48    49    50
[2,]  1521  1600  1681  1764  1849  1936  2025  2116  2209  2304  2401  2500
      [,51] [,52] [,53] [,54] [,55] [,56] [,57] [,58] [,59] [,60] [,61] [,62]
[1,]    51    52    53    54    55    56    57    58    59    60    61    62
[2,]  2601  2704  2809  2916  3025  3136  3249  3364  3481  3600  3721  3844
      [,63] [,64] [,65] [,66] [,67] [,68] [,69] [,70] [,71] [,72] [,73] [,74]
[1,]    63    64    65    66    67    68    69    70    71    72    73    74
[2,]  3969  4096  4225  4356  4489  4624  4761  4900  5041  5184  5329  5476
      [,75] [,76] [,77] [,78] [,79] [,80] [,81] [,82] [,83] [,84] [,85] [,86]
[1,]    75    76    77    78    79    80    81    82    83    84    85    86
[2,]  5625  5776  5929  6084  6241  6400  6561  6724  6889  7056  7225  7396
      [,87] [,88] [,89] [,90] [,91] [,92] [,93] [,94] [,95] [,96] [,97] [,98]
[1,]    87    88    89    90    91    92    93    94    95    96    97    98
[2,]  7569  7744  7921  8100  8281  8464  8649  8836  9025  9216  9409  9604
      [,99] [,100]
[1,]    99   100
[2,]  9801 10000
```

...

Tip

Recall the discussion in class on how R fills in matrices

In the next part, we will discover how knowledge of the way in which a matrix is stored in memory can inform better code choices. To this end, the following function takes an input n and creates an $n \times n$ matrix with random entries.

```

generate_matrix <- function(n){
  return(
    matrix(
      rnorm(n^2),
      nrow=n
    )
  )
}

```

For example:

```
generate_matrix(4)
```

```

      [,1]      [,2]      [,3]      [,4]
[1,]  0.1217731  0.1529565 -0.5807544  0.7608220
[2,] -0.3128948 -0.2352783  1.5194928 -0.7895756
[3,]  0.8409010  0.8757768  0.5002945  0.6745095
[4,]  1.0423601  1.0226644 -1.8205939  1.2358449

```

Let M be a fixed 50×50 matrix

```

M <- generate_matrix(50)
mean(M)

```

```
[1] -0.03350234
```

2. Write a function `row_wise_scan` which scans the entries of M one row after another and outputs the number of elements whose value is ≥ 0 . You can use the following **starter code**

```

row_wise_scan <- function(x){
  n <- nrow(x)
  m <- ncol(x)

  # Insert your code here
  count <- 0
  for(i in 1:n){
    for(j in 1:m){
      if(x[i,j] >= 0){
        count <- count + 1
      }
    }
  }
}

```

```

    }
  }

  return(count)
}

```

3. Similarly, write a function `col_wise_scan` which does exactly the same thing but scans the entries of `M` one column after another

```

col_wise_scan <- function(x){
  count <- 0

  n <- nrow(x)
  m <- ncol(x)

  for(i in 1:m){
    for(j in 1:n){
      if(x[j,i] >= 0){
        count <- count + 1
      }
    }
  }

  return(count)
}

```

You can check if your code is doing what it's supposed to using the function [here](#)¹

- Between `col_wise_scan` and `row_wise_scan`, which function do you expect to take shorter to run? Why?
- Write a function `time_scan` which takes in a method `f` and a matrix `M` and outputs the amount of time taken to run `f(M)`

¹If your code is right, the following code should evaluate to be `TRUE`

```

sapply(1:100, function(i) {
  x <- generate_matrix(100)
  row_wise_scan(x) == col_wise_scan(x)
}) %>% sum == 100

```

```
time_scan <- function(f, M){
  initial_time <- Sys.time() # Write your code here
  f(M)
  final_time <- Sys.time() # Write your code here

  total_time_taken <- final_time - initial_time
  return(total_time_taken)
}
```

Provide your output to

```
list(
  row_wise_time = time_scan(row_wise_scan, M),
  col_wise_time = time_scan(col_wise_scan, M)
)
```

Which took longer to run?

6. Repeat this experiment now when:

- M is a 100×100 matrix
- M is a 1000×1000 matrix
- M is a 5000×5000 matrix

What can you conclude?

Appendix

Print your R session information using the following command

```
sessionInfo()
```

```
R version 4.2.2 (2022-10-31 ucrt)
Platform: x86_64-w64-mingw32/x64 (64-bit)
Running under: Windows 10 x64 (build 22621)
```

```
Matrix products: default
```

```
locale:
```



```
[1] LC_COLLATE=English_United States.utf8
[2] LC_CTYPE=English_United States.utf8
[3] LC_MONETARY=English_United States.utf8
[4] LC_NUMERIC=C
[5] LC_TIME=English_United States.utf8
```

attached base packages:

```
[1] stats      graphics  grDevices datasets  utils      methods    base
```

loaded via a namespace (and not attached):

```
[1] compiler_4.2.2 fastmap_1.1.0 cli_3.6.0      htmltools_0.5.4
[5] tools_4.2.2    yaml_2.3.7    rmarkdown_2.20 knitr_1.42
[9] xfun_0.36      digest_0.6.31 jsonlite_1.8.4 rlang_1.0.6
[13] renv_0.16.0-53 evaluate_0.20
```