# Homework 1

# Ziyao Yang

## Table of contents

| Question 2 |
|------------|
| Question 2 |
|            |

Link to the Github repository

Due: Sun, Jan 29, 2023 @ 11:59pm

Please read the instructions carefully before submitting your assignment.

- 1. This assignment requires you to:
  - Upload your Quarto markdown files to a git repository
  - Upload a PDF file on Canvas
- 2. Don't collapse any code cells before submitting.
- 3. Remember to make sure all your code output is rendered properly before uploading your submission.

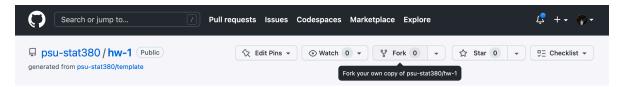
Please add your name to the the author information in the frontmatter before submitting your assignment.

#### Question 1



In this question, we will walk through the process of *forking* a git repository and submitting a *pull request*.

1. Navigate to the Github repository here and fork it by clicking on the icon in the top right



Provide a sensible name for your forked repository when prompted.

2. Clone your Github repository on your local machine

```
$ git clone <<insert your repository url here>>
$ cd hw-1
```

Alternatively, you can use Github codespaces to get started from your repository directly.

3. In order to activate the R environment for the homework, make sure you have renv installed beforehand. To activate the renv environment for this assignment, open an instance of the R console from within the directory and type

```
renv::activate()
```

Follow the instrutions in order to make sure that renv is configured correctly.

- 4. Work on the reminaing part of this assignment as a .qmd file.
  - Create a PDF and HTML file for your output by modifying the YAML frontmatter for the Quarto .qmd document
- 5. When you're done working on your assignment, push the changes to your github repository.

6. Navigate to the original Github repository here and submit a pull request linking to your repository.

Remember to **include your name** in the pull request information!

If you're stuck at any step along the way, you can refer to the official Github docs here <sup>1</sup>

## Question 2



Consider the following vector

```
my_vec <- c(
    "+0.07",
    "-0.07",
    "+0.25",
    "-0.84",
    "+0.32",
    "-0.24",
    "-0.97",
    "-0.36",
    "+1.76",
    "-0.36")
```

For the following questions, provide your answers in a code cell.

1. What data type does the vector contain?

```
typeof(my_vec)
```

#### [1] "character"

2. Create two new vectors called my\_vec\_double and my\_vec\_int which converts my\_vec to Double & Integer types, respectively,

<sup>&</sup>lt;sup>1</sup>I was unable to render the pdf file. I upgrade R version and all the packages, followed the instructions returned by "renv::status()" and synchronized the project with the lockfile. I'm not sure if these are the right process to solve the problem but I did successfully rendered pdf after these steps.

```
my_vec_double <- as.double(my_vec)
my_vec_double

[1] 0.07 -0.07 0.25 -0.84 0.32 -0.24 -0.97 -0.36 1.76 -0.36

my_vec_int <- as.integer(my_vec)
my_vec_int

[1] 0 0 0 0 0 0 0 0 1 0</pre>
```

- 3. Create a new vector my\_vec\_bool which comprises of:
  - TRUE if an element in  $my_vec_double$  is  $\leq 0$
  - FALSE if an element in  $my\_vec\_double$  is  $\geq 0$

```
my_vec_bool <- ifelse(my_vec_double<=0, yes = TRUE, no = FALSE)
# ifelse(my_vec_double<=0, TRUE, FALSE)
# my_vec_bool <- my_vec_double<=0
my_vec_bool</pre>
```

[1] FALSE TRUE FALSE TRUE FALSE TRUE TRUE TRUE FALSE TRUE

How many elements of my\_vec\_double are greater than zero?

```
sum(my_vec_double > 0)
```

[1] 4

4. Sort the values of my\_vec\_double in ascending order.

```
sort(my_vec_double)
```

[1] -0.97 -0.84 -0.36 -0.36 -0.24 -0.07 0.07 0.25 0.32 1.76

## Question 3

• 50 points

In this question we will get a better understanding of how R handles large data structures in memory.

1. Provide R code to construct the following matrices:

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} \quad \text{and} \quad \begin{bmatrix} 1 & 2 & 3 & 4 & 5 & \dots & 100 \\ 1 & 4 & 9 & 16 & 25 & \dots & 10000 \end{bmatrix}$$



Recall the discussion in class on how R fills in matrices

In the next part, we will discover how knowledge of the way in which a matrix is stored in memory can inform better code choices. To this end, the following function takes an input nand creates an  $n \times n$  matrix with random entries.

```
generate_matrix <- function(n){</pre>
    return(
         matrix(
             rnorm(n^2),
             nrow=n
         )
    )
}
```

For example:

```
generate_matrix(4)
```

```
[,1]
                                  [,3]
                      [,2]
                            0.47027687 -0.3500899
     1.5318180 -0.9983621
[2,] -0.2922491
                 0.4676499
                            1.54509368 0.5047034
                1.7944968 -0.25521193 -0.5990463
[3,]
     0.3214969
     0.6873778
                1.6973251 -0.01395378 -0.9796172
```

Let M be a fixed  $50 \times 50$  matrix

```
M <- generate_matrix(50)</pre>
mean(M)
```

[1] 0.01430765

2. Write a function row\_wise\_scan which scans the entries of M one row after another and outputs the number of elements whose value is  $\geq 0$ . You can use the following starter code

```
row_wise_scan <- function(x){
    n <- nrow(x)
    m <- ncol(x)

# Insert your code here
    count <- 0
    for(...){
        if(...){
            count <- count + 1
            }
        }
    }
    return(count)
}</pre>
```

3. Similarly, write a function col\_wise\_scan which does exactly the same thing but scans the entries of M one column after another

```
col_wise_scan <- function(x){
    count <- 0

    ... # Insert your code here
    return(count)
}</pre>
```

You can check if your code is doing what it's supposed to using the function here<sup>2</sup>

4. Between col\_wise\_scan and row\_wise\_scan, which function do you expect to take shorter to run? Why?

```
sapply(1:100, function(i) {
    x <- generate_matrix(100)
    row_wise_scan(x) == col_wise_scan(x)
}) %>% sum == 100
```

<sup>&</sup>lt;sup>2</sup>If your code is right, the following code should evaluate to be TRUE

5. Write a function time\_scan which takes in a method f and a matrix M and outputs the amount of time taken to run f(M)

```
time_scan <- function(f, M){
   initial_time <- ... # Write your code here
   f(M)
   final_time <- ... # Write your code here

   total_time_taken <- final_time - initial_time
   return(total_time_taken)
}</pre>
```

Provide your output to

```
list(
    row_wise_time = time_scan(row_wise_scan, M),
    col_wise_time = time_scan(row_wise_scan, M)
)
```

Which took longer to run?

- 6. Repeat this experiment now when:
  - M is a  $100 \times 100$  matrix
  - M is a  $1000 \times 1000$  matrix
  - M is a  $5000 \times 5000$  matrix

What can you conclude?

# **Appendix**

Print your R session information using the following command

```
R version 4.2.2 (2022-10-31)
Platform: aarch64-apple-darwin20 (64-bit)
Running under: macOS Monterey 12.6

Matrix products: default
```

BLAS: /Library/Frameworks/R.framework/Versions/4.2-arm64/Resources/lib/libRblas.0.dylib LAPACK: /Library/Frameworks/R.framework/Versions/4.2-arm64/Resources/lib/libRlapack.dylib

#### locale:

[1] en\_US.UTF-8/en\_US.UTF-8/en\_US.UTF-8/C/en\_US.UTF-8/en\_US.UTF-8

### attached base packages:

[1] stats graphics grDevices datasets utils methods base

### loaded via a namespace (and not attached):

| [1] compiler_4.2.2 | $fastmap_1.1.0$ | cli_3.6.0 | htmltools_0.5.4 |
|--------------------|-----------------|-----------|-----------------|
|--------------------|-----------------|-----------|-----------------|

[5] tools\_4.2.2 rstudioapi\_0.14 yaml\_2.3.7 rmarkdown\_2.20

[9] knitr\_1.42 xfun\_0.36 digest\_0.6.31 jsonlite\_1.8.4

[13] rlang\_1.0.6 renv\_0.16.0-53 evaluate\_0.20